

The Linker



The Linker

- Links object files into an executable or library
- Resolves **symbols** from different sources
- Fixes **addresses** for functions & variables
- Produces **executables, DLLs (Windows), SOs (Linux)**

Static Libraries

- $+$
- $*$
- $^$





Full Source Compilation



add → mul → pow → main



```
g++ add.cpp mul.cpp pow.cpp main.cpp
```

-  All `.cpp` files compiled and linked in one step
-  No dependency order needed

Object File Linking

Build objects first, then link



```
g++ add.cpp -c  
...  
g++ add.o mul.o pow.o main.cpp
```

-  Compile each source separately
-  Still no link order issues

Static Library Linking

 Libraries need correct order!

```
g++ add.cpp -c  
ar rcs libadd.a add.o  
...  
  
g++ main.cpp libpow.a libmul.a libadd.a
```

-  Linker reads **left to right**
-  Put libraries **after** the code that needs them



Inspecting with `nm`

```
nm libmul.a -C
```

```
mul.o:  
                U _GLOBAL_OFFSET_TABLE_  
                U add(int, int)  
000000000000000000 T mul(int, int)
```

Link Stages: Windows

1. Object files (`.obj`) compiled
2. Linker resolves **all symbols**
3. Statically links `.lib` files
4. Produces `.exe` or `.dll`

Link Stages: Linux

1. Object files (`.o`) compiled
2. Linker resolves symbols **in order**
3. Links `.a` and `.so` files
4. Produces `executable`

⚠ **Order matters in Linux!**

```
g++ main.o -ladd -lmul -lpow # ❌ Undefined reference  
g++ main.o -lpow -lmul -ladd # ✅ Correct order
```

Undefined Reference

- 🔍 Use `nm`, `grep` to find where it's defined
- 🧱 Check link command (`link.txt` in CMake)
- + Ensure dependency is **after** its user
- 🛠️ Fix link order or use `LINK_GROUP`

```
find . -type f -name '*.a' -exec nm -C {} + 2>/dev/null
```



Undefined Symbol



Missing Library or Flag

`pthread_create`, `pthread_join`

`-lpthread`

`boost::filesystem::path`

`-lboost_filesystem`

`cv::Mat`

`-lopencv_core`

`dlopen`, `dlsym`, `dlclose`

`-ldl`

`std::thread`, `std::mutex`

`-pthread`

`zlibVersion`, `inflate`, `deflate`

`-lz`

`curl_easy_init`, `curl_easy_setopt`

`-lcurl`

`glBegin`, `glVertex3f`, `glEnd`

`-lGL` (OpenGL)

`glutInit`, `glutCreateWindow`

`-lglut`



Which libraries will be linked ?

Linux Flags

- `-L<dir>` — Add search path
- `-l<name>` — Link `lib<name>.so` / `.a`
- `-l:filename.a` — Exact archive

Windows (MSVC)

- Linker -> Input -> Additional Dependencies
- `#pragma comment(lib, "mylib.lib")` in code



CMake: Adding Libraries

- `target_link_libraries()` connects targets

```
add_library(mylib STATIC file.cpp)
target_link_libraries(app PRIVATE mylib)
```



Dynamic Libraries

- +
- *
- ^

	🧱 Static (<code>.a</code> , <code>.lib</code>)	🔗 Dynamic (<code>.so</code> , <code>.dll</code>)
Linked	At compile/link time	At runtime
Included in	Final executable	External file
Size	Larger executable	Smaller executable
Flexibility	Less (update=rebuild)	More (swap <code>.dll</code> / <code>.so</code>)
Dependencies	None at runtime	<code>.so</code> / <code>.dll</code> must be present

💡 Use **static** if you can.



Create a DLL (Windows)

-  Change `CMakeLists.txt`:

```
add_library(add SHARED add.cpp)
```

-  Mark functions for export:

```
__declspec(dllexport)  
int add(int a, int b) {  
    return a + b;  
}
```





Portable Export Macro

```
#ifdef MYLIB_EXPORTS
    #define MYLIB_API __declspec(dllexport)
#else
    #define MYLIB_API __declspec(dllimport)
#endif
```

```
MYLIB_API int add(int a, int b);
```

```
target_compile_definitions(mylib PRIVATE MYLIB_EXPORTS)
```

Linking with a DLL (Windows)

-  Link against `.lib` (import library)
-  Loads `.dll` at runtime

DLL Search Order:

1. Executable directory
2. Current directory
3. System `PATH`



Create a `.so` (Linux)


-  Change `CMakeLists.txt`:

```
add_library(add SHARED add.cpp)
```

-  No need for `__declspec(dllexport)`

```
int add(int a, int b) {  
    return a + b;  
}
```

Linking with a `.so` (Linux)

-  Link with `-ladd`
- Provide `.so` at runtime

`.so` Search Order:

1. `DT_RPATH` (Deprecated)
2. `LD_LIBRARY_PATH`
3. `DT_RUNPATH`
4. `/etc/ld.so.cache`

Link Stages: Linux

✅ `.so` files can have **undefined symbols**

⚠️ They're checked **only when linking the final executable**

The `--no-undefined` flag changes that:

```
add_library(pow SHARED pow.cpp)

set_target_properties(pow PROPERTIES
    LINK_FLAGS "-Wl,--no-undefined"
)
target_link_libraries(pow PRIVATE mul add)
```

CMake

By default :

- CMake **adds** run path to the executable
- CMake **removes** the run path when installing

Recommended :

```
set(CMAKE_INSTALL_PREFIX ${CMAKE_BINARY_DIR}/install)

set_target_properties(main PROPERTIES
  INSTALL_RPATH "$ORIGIN"
  SKIP_BUILD_RPATH ON
)
```



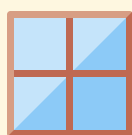

CMake Global vs Project

- Global

```
set(CMAKE_SKIP_BUILD_RPATH TRUE)
```

- More fine-grained control

```
set_target_properties(main  
  PROPERTIES  
  SKIP_BUILD_RPATH TRUE  
)
```



Shared Library Differences

Feature	Linux (<code>.so</code>)	Windows (<code>.dll</code>)
Link at <code>.so</code> build	✗ (lazy)	✓ (must resolve)
Search order	<code>LD_LIBRARY_PATH</code> , <code>rpath</code> , <code>ld.so.cache</code>	Current dir, PATH, system dirs
Default visibility	Public	Hidden (needs <code>__declspec(dllexport)</code>)
Extension	<code>.so</code>	<code>.dll</code> + <code>.lib</code> (import)

Tool



Dependency walker

- <https://www.dependencywalker.com/>
- <https://github.com/lucasg/Dependencies>

Runtime Linker Debugging Tools

✓ **Linux:** `LD_DEBUG`, `LD_PRELOAD`

✓ **Windows:** `procmon`





LD_DEBUG

- ◆ Debug **dynamic linker activity** 🏗️
- ◆ Show **symbol resolution, library loading** 🕵️

```
LD_DEBUG=all ./my_program # Show everything 👁️  
LD_DEBUG=libs ./my_program # Library loading 🔍  
LD_DEBUG=symbols ./my_program # Symbol lookup 📖
```

DO NOT

Exporting All Symbols

- ◆ By default, all symbols may be exported 
- ◆ **Problem:** Exposing unnecessary functions may cause **symbol conflicts** 
- ◆ **Example:** Exporting 3rd-party libraries like **Boost, IPP, OpenCV**



What Can Go Wrong?

- ✗ **Symbol Conflicts** – Runtime crashes, wrong ABI ⚡
- ✗ **Linking Issues** – Different version might be used across libraries
📌
- ✗ **Unintended ABI Exposure** – Internal functions can accidentally be used 🔧

How to Verify Symbol Exports

✓ Linux: Use `nm` or `objdump`

```
nm -D myLib.so  # Lists exported symbols  
objdump -T myLib.so  # Shows dynamic symbols
```

◆ Check exports before releasing shared libraries! 

Function Visibility

✓ **Windows (DLLs)** → Uses `__declspec(dllexport)` & `__declspec(dllimport)`.

✓ **Linux (SOs)** → Uses `__attribute__((visibility("default")))`.

Notes :

- `__declspec(dllimport)` is optional, for functions.
- Linux - default is visible, Windows - default is invisible.
- Which functions are exported: `nm -D` , `objdump` / `dependencies`

◆ Function Visibility in Windows

```
#ifdef BUILD_DLL
#define API_EXPORT __declspec(dllexport)
#else
#define API_EXPORT __declspec(dllimport)
#endif

API_EXPORT void myFunction(); // Exported function
```

✓ **Importing (from EXE or another DLL) → Use**



```
__declspec(dllimport)
```



Function Visibility in Linux I

```
#define MYLIB_API __attribute__((visibility("default")))  
  
MYLIB_API int add(int a, int b);
```

```
target_compile_options(mylib PRIVATE -fvisibility=hidden)
```

-  Hides symbols by default
-  Exports only tagged symbols (like `MYLIB_API`)



Function Visibility in Linux II

-  Create a version script (`exports.map`):

```
{  
  global:  
    add;  
  local:  
    *;  
};
```







-  Use it in CMake:

```
set_target_properties(add PROPERTIES  
  LINK_FLAGS "-Wl,--version-script=${CMAKE_SOURCE_DIR}/exports.map"  
)
```





Advanced

- Manual loading
- Name mangling
- Versioning (`so.1.2.3`)

Manually Loading DLLs

-  Load at runtime (plugins, late binding)
-  Use `dlsym` / `GetProcAddress` / `boost::dll`
-  Flexible, optional dependencies
-  Must know symbol names
-  `extern "C"` avoids name mangling
-  Manage global/static init

Manual Loading: Pitfalls

-  Hard to debug missing or broken symbols
-  Symbol names must match exactly
-  No linker error → runtime crash
-  Test for load failure (`nullptr`, try/catch)

Boost::dll Example

```
// Load the shared library
#ifdef _WIN32
boost::dll::shared_library lib("mylib.dll");
#else
boost::dll::shared_library lib("libmylib.so");
#endif
// Import the function dynamically
auto hello_func = lib.get<void()>("hello");
// Call the function
hello_func();
```

Name Mangling

- ◆ C++ compilers modify function names
- ◆ Compiler specific
- ◆ C does NOT mangle names

```
void foo();    // Regular function
void foo(int); // Overloaded function
namespace A {
    void foo(); // Namespaced function
}
```

Name Mangling in GCC/Clang (nm)

```
_Z3foov      # void foo()  
_Z3fooRi     # void foo(int)  
_N1A3fooEv   # void A::foo()
```

Name Mangling in MSVC (dumpbin)

```
?foo@@YAXXZ  # void foo()  
?foo@@YAXH@Z # void foo(int)  
?foo@A@@YAXXZ # void A::foo()
```

extern "C"

```
extern "C" void foo(); // No name mangling applied
extern "C" {
    void bar();
    int baz(int);
}
```

✓ Use `extern "C"` for

- C++ code that needs C-compatible APIs
- Dynamic libraries that must be callable from C
- Avoiding cross-compiler mangling issues

DLL Hell (Windows)




⚠ Version conflicts

- Different apps require different versions of the same DLL
- **Example:** `gda1.dll`
- **Fix:**
 - Use local DLLs
 - Rename
 - [**SxS** (Side-by-Side assemblies)]

SO Versioning in Linux

◆ Structure:

```
libmylib.so → libmylib.so.1 → libmylib.so.1.2.3
```

- **1** = Major (Breaking changes )
- **2** = Minor (New features )
- **3** = Patch (Bug fixes )

SO Version in CMake

```
add_library(mylib SHARED mylib.cpp)

set_target_properties(mylib PROPERTIES
    VERSION 1.2.3      # Full version
    SOVERSION 1        # Major version
)
```

◆ Creates:

```
libmylib.so → libmylib.so.1 → libmylib.so.1.2.3
```




LD_PRELOAD

- ◆ Inject shared libraries 
- ◆ Override functions without rebuilding 

```
LD_PRELOAD=/path/to/mylib.so ./my_program
```

Common linker problems, its cause, how to debug, and solutions.



-fPIC

Problem

```
/usr/bin/ld: mylib.a(myfile.o): relocation R_X86_64_32S against `'.text'
can not be used when making a shared object; recompile with -fPIC
```

Solution



Compile with **-fPIC**



MSVC

#pragma detect_mismatch

Key	Used For	Example Value
<code>_MSC_VER</code>	Compiler version	<code>"1934"</code>
<code>RuntimeLibrary</code>	Static vs. dynamic CRT	<code>"MD_DynamicDebug"</code>
<code>_ITERATOR_DEBUG_LEVEL</code>	STL iterator checks level	<code>"0"</code> / <code>"2"</code>

⚠ Mismatch = link error (LNK2038)

✗ `_ITERATOR_DEBUG_LEVEL` Mismatch (MSVC)

Problem

```
error LNK2038: mismatch detected for '_ITERATOR_DEBUG_LEVEL':  
value '0' doesn't match value '2'
```

Solution

✓ Ensure **consistent build types**

✗ MDd_DynamicDebug Mismatch (MSVC)

Problem

```
error LNK2038: mismatch detected for 'RuntimeLibrary':  
value 'MTd_StaticDebug' doesn't match value 'MDd_DynamicDebug' in program.obj  
C:\Data\ip_core.lib(ip_core.obj)
```

✗ Cannot find `mydll.lib` (MSVC)

Problem

```
LINK : fatal error LNK1104: cannot open file 'Debug\add.lib'
```

Solution

✓ In windows, a `.lib` file is created **Only if there are exported symbols**. Check `__declspec(dllexport)`

❌ Incorrect Shared Object (.so) Link Order

Problem

```
/usr/bin/ld: undefined reference to `myFunction`
```

Solution

✅ **Correct order**

✅ **Use grouping:**

```
g++ -Wl,--start-group -lfoo -lbar -Wl,--end-group -o myapp
```


✖ Missing `.so` File at Runtime

Problem

```
./myapp: error while loading shared libraries: libmylib.so: cannot open shared object file
```

Solution

- ✓ Check if the library is found
- ✓ Fix missing paths `LD_LIBRARY_PATH`

✗ Wrong shared object loaded at Runtime

Solution

✓ Check which DLL is loaded

- Windows : Debug -> Windows -> Modules
- Linux : gdb -> `info shared`



? Q1: What does the linker do?

1. Runs your code
2. Edits the source files
3. Resolves symbols and produces an executable
4. Downloads libraries

? Q2: What happens if you link a static library in the wrong order?

1. Nothing, order doesn't matter
2. The build will be faster
3. You get undefined reference errors
4. All functions get linked anyway

? Q3: What does `-lfoo` link against?

1. `foo.c`
2. `libfoo.so` or `libfoo.a`
3. `foo.cpp`
4. `foo.o`

? Q4: What does `-L` do?

1. Links a library
2. Specifies a source directory
3. Adds a library search path
4. Loads a DLL

? Q5: Which tool shows exported symbols in a Linux `.so`?

1. nm -D
2. make
3. grep
4. g++

? Q6: What pragma is used in MSVC to catch mismatched settings?

1. `#pragma mismatch`
2. `#pragma error`
3. `#pragma detect_mismatch`
4. `#pragma validate`

? Q7: Which visibility is default on Windows for DLLs?

1. Visible
2. Hidden
3. Public
4. External

? Q8: What is the purpose of `__declspec(dllexport)` ?

1. Export symbols
2. Load a DLL
3. Import symbols from DLL
4. Hide internal functions

? Q9: What happens when `LD_PRELOAD` is used?

1. Loads libraries after program starts
2. Overrides linked symbols
3. Speeds up linking
4. Changes runtime paths

? Q10: Which variable controls runtime search path for `.so`?

1. `PATH`
2. `LD_RUN_PATH`
3. `LD_LIBRARY_PATH`
4. `LD_BIN`

? Q11: What does `target_link_libraries()` do in CMake?

1. Adds include path
2. Creates shared library
3. Links targets to dependencies
4. Starts the build process

? Q12: What does `--no-undefined` flag do?

1. Skips symbols
2. Allows missing functions
3. Requires all symbols to be resolved
4. Only works in debug

? Q13: What does

`add_library(foo SHARED foo.cpp)` do?

1. Creates a static library
2. Creates a header file
3. Creates a shared object (.so/.dll)
4. Runs `foo.cpp`

? Q14: How to hide symbols by default in Linux?

1. Use `-fPIC`

2. Use `-O3`

3. Use `-fvisibility=hidden`

4. Use `strip`

? Q15: Which tool can inspect `.lib` or `.obj` files on Windows?

1. ldd
2. dumpbin
3. objdump
4. strings