

# C++ Quiz Challenge



## Join the Game!

<https://play.blooket.com/>

40+ Questions • Git • Memory • Multithreading • C++17 and more!

```
$ git stash
```

## Q1. 🐙 What does 'git stash' do?

- A) 🗑️ Deletes untracked files
- B) 📦 Commits all staged changes
- C) 🔍 Resets to last commit
- D) 💾 Saves uncommitted changes, cleans working directory

## Q2. What is the state of myfile.cpp after these two commands?

```
$ git add myfile.cpp  
$ git reset HEAD -- myfile.cpp
```

- A)  Committed to branch
- B)  Deleted from repo
- C)  Modified, unstaged
- D)  Staged, ready to commit

## Q3. ⚙ What does 'my\_rules.csv binary' in .gitattributes do?

```
# .gitattributes  
my_rules.csv  binary
```

- A)  Forces CRLF on checkout
- B)  Compresses the file
- C)  Ignores the file in commits
- D)  git won't try to merge it

# Git

---

Version control: init, add, commit, push

Branching: one branch = one task

Stash: save work without committing

Pre-commit hooks: automate checks

CRLF: always use LF (eol=lf)

CI/CD: compile on isolated machines

## Q5. 🔐 What is the correct order of these 4 steps?

```
$ g++ -E simple.cpp -o simple.i  
$ g++ -S simple.i -o simple.s  
$ g++ -c simple.s -o simple.o  
$ g++ simple.o -o simple
```

- A) Compile → Preprocess → Assemble → Link
- B) Assemble → Compile → Preprocess → Link
- C) Preprocess → Compile → Assemble → Link
- D) Preprocess → Assemble → Compile → Link

## Q6. 🔎 What does 'ldd simple' show?

```
$ ldd simple
 libstdc++.so.6 => /usr/lib/...
 libc.so.6        => /lib/...
 linux-vdso.so.1  (0x...)
```

- A) 🏷️ All symbols defined in the binary
- B) 💻 System calls used at runtime
- C) 📁 Static libraries linked into the binary
- D) 📚 Shared libraries the binary depends on

```
# CMakeLists.txt
set(CMAKE_SYSTEM_NAME Linux)
set(CMAKE_SYSTEM_PROCESSOR aarch64)
set(CMAKE_CXX_COMPILER
    aarch64-linux-gnu-g++)
```

## Q7. What does **CMAKE\_SYSTEM\_NAME** set in a CMake toolchain file?

- A)  The host machine's OS name
- B)  The build system generator
- C)  The compiler vendor name
- D)  The target OS for cross-compilation

## ☐ Toolchain & Build

---

Toolchain: compiler, linker, assembler, debugger

Compilation: Preprocess -> Compile -> Assemble -> Link

ldd: inspect shared library dependencies

glibc: tightly coupled to Linux kernel

Static vs dynamic linking trade-offs

Cross-compilation: CMAKE\_SYSTEM\_NAME / PROCESSOR

```
// myheader.h  
#pragma once  
  
int foo();
```

## Q9. What does '#pragma once' do in a header file?

- A)  Disables all warnings in the file
- B)  Marks the file as system header
- C)  Prevents the file from being included more than once
- D)  Forces the file to compile before all others

## **Q10. This code fails to compile on Windows. What is the cause?**

```
#include <algorithm>
#include <Windows.h>

int k = std::min(3,4);
```

- A)**  algorithm is not available on Windows
- B)**  Missing #include for <windows.h>
- C)**  Windows.h defines min/max as macros
- D)**  std::min needs two identical types

```
#define SQR(x)  x*x  
  
std::cout << SQR(4+1);
```

**Q11.** 🤯 **#define SQR(x)**  
**x\*x** What is the output of  
**SQR(4+1)?**

- A) 9
- B) 21
- C) 25
- D) 1

## □ Preprocessor

---

#include: text substitution, watch search dirs  
#pragma once: prevent multiple inclusion  
#define: avoid macros, prefer constexpr  
Windows.h: define NOMINMAX before including  
g++ -E: dump preprocessed output to debug  
clang-tidy: detect include issues

```
uint64_t foo() {  
    const int N = 1024;  
    uint64_t total = 0;  
    for(int i=0;i<N;++i)  
        total += i*i;  
    return total;  
} // g++ -O3
```

## Q13. 🚀 What does g++ -O3 return for this function?

- A) ⚡ A compiler error: loop is infinite
- B) 0 Zero, loop body is optimized away
- C) 📊 A single constant (constant folding)
- D) ⚡ The loop runs 1024 times at runtime

```
#include <climits>
int main() {
    int max = INT_MAX;
    max++; // ???
}
```

## Q14. 💀 What is the behavior of 'max++' when max == INT\_MAX?

- A) ⚡ Compiler error: overflow detected
- B) 🔐 Wraps to INT\_MIN (guaranteed)
- C) 0 Result is 0
- D) 💣 Undefined behavior (signed overflow)

```
$ g++ -fsanitize=address \  
      -fsanitize=undefined \  
      main.cpp
```

## Q15. 🩺 What do sanitizer flags like - fsanitize=address do?

- A) ✂️ Strip debug symbols from the binary
- B) 🔔 Disable compiler optimizations
- C) ⚡ Speed up execution by removing bounds checks
- D) 🔎 Add runtime checks that detect memory errors

## □ Compiler

---

-O0/-O2/-O3: optimization levels

Constant folding: loops -> single instruction

Undefined behavior: silent data corruption

Sanitizers: -fsanitize=address,undefined

-Wall -Wextra -Werror: enable all warnings

clang vs MSVC: error message quality

## Q17. Which link command is correct on Linux?

```
# Which is correct on Linux?  
g++ main.o -lpow -lmul -ladd  
g++ main.o -ladd -lmul -lpow
```

- A)  Use --start-group for all cases
- B) g++ main.o -ladd -Imul -Ipow
- C) g++ main.o -Ipow -Imul -ladd
- D)  Order does not matter on Linux

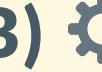
## Q18. 🔎 In nm output, what does 'T' mean next to a symbol?

```
$ nm -C libmul.a
U add(int, int)
T mul(int, int)
```

- A)  In a data section
- B)  Template function
- C)  Undefined, must be linked
- D)  Defined in the text (code)  
section

```
// add.cpp  (Windows DLL)
int add(int a, int b) {
    return a + b;
}
```

## Q19. A Windows DLL is built but no .lib file is created. Why?

- A)  No `__declspec(dllexport)` symbols in the code
- B)  CMake forgot to set SHARED keyword
- C)  Wrong architecture (x86 vs x64)
- D)  Missing /MD compiler flag

# □ Linker

---

Links .o files into executable / .so / .dll

Linux: library order matters (left to right)

nm: inspect symbols (T=defined, U=undefined)

Static .a vs dynamic .so/.dll trade-offs

Windows: \_\_declspec(dllexport) to export

Linux: -fvisibility=hidden + \_\_attribute\_\_

```
#include <iostream>
int main() {
    std::cout << sizeof(long);
}
// 64-bit Linux vs Windows?
```

## Q21. **sizeof(long)** on 64-bit Linux vs 64-bit Windows?

- A) Both: 4
- B) Linux: 8, Windows: 4
- C) Both: 8
- D)  Linux: 4, Windows: 8

```
char c = 200;  
if (c > 0)  
    std::cout << "Positive";  
else  
    std::cout << "Negative";  
// x86 GCC vs ARM GCC?
```

## Q22. 🤔 **char c = 200 – what is the output on x86 GCC vs ARM GCC?**

- A) – Both print Negative
- B) 💀 Undefined behavior on both platforms
- C) + Both print Positive
- D) ✎ Different – char signedness is implementation-defined

```
#pragma pack(push, 1)
struct S {
    char a;    // 1 byte
    int b;     // 4 bytes
    short c;   // 2 bytes
};
#pragma pack(pop)
```

**Q23.**  **What is sizeof(S)  
after #pragma  
pack(push,1)?**

- A)** 12
- B)** 4
- C)** 7
- D)** 8

## □ ABI

---

ABI = Application Binary Interface

Rule 1: Use POD types only

Rule 2: Use fixed-width types (`int32_t...`)

Rule 3: `#pragma pack(push,1)` for structs

Rule 4: Export only needed symbols

Rule 5: `extern "C"` to prevent name mangling

## Q25. What does 'strip app' do to the binary?

```
$ g++ -g -O2 main.cpp -o app  
$ strip app
```

- A)  Deletes the binary from disk
- B)  Encrypts the binary for distribution
- C)  Removes optimizations, adds debug info
- D)  Removes debug symbols, shrinks binary size

## Q26. What does 'watch var' do in gdb?

```
(gdb) watch var  
(gdb) rwatch var  
(gdb) awatch var
```

- A)  Stops execution when var is written
- B)  Stops execution when var is read
- C)  Prints var value every step
- D)  Sets var to zero and continues

## Q27. ⚡ What do these two commands enable on Linux?

```
$ ulimit -c unlimited  
$ echo "/tmp/core.%e.%p" |  
  sudo tee /proc/sys/kernel/core_pattern
```

- A)  Core dump creation when a process crashes
- B)  Debug symbol loading from /proc
- C)  Remote gdbserver on port 2000
- D)  Full memory logging to /tmp

## □ Debugging

---

Debug info: -g (Linux DWARF) / /Zi (Windows PDB)

RelWithDebInfo: -O2 + -g (best of both worlds)

strip: remove debug info from binary

gdb: break, next, step, bt, watch, p

Core dumps: ulimit -c unlimited

Remote: gdbserver :2000 ./app

```
void foo(int a) {  
    int x = 42;  
    bar(x);  
}
```

**Q29. 🧱 What memory region holds local variables like 'int x'?**

- A)  Stack
- B)  ROM
- C)  Global/Static
- D)  Heap

```
int& get_ref() {  
    int x = 42;  
    return x;  
}  
  
int main() {  
    int& r = get_ref();  
    return r;  
}
```

## Q30. 💀 What is wrong with this get\_ref() function?

- A) 📦 x is allocated on the heap
- B) 🔁 Infinite recursion
- C) 🚫 Missing return type
- D) 🔥 Returns a dangling  
reference to a local variable

```
void leak() {  
    int* p = new int[100];  
    if (!OpenFile())  
        return;  
    delete[] p;  
}
```

## Q31. 🍹 What bug does this leak() function contain?

- A) 📈 Buffer overflow
- B) 💀 Use-after-free
- C) 💧 Memory leak when OpenFile() returns false
- D) 🔥 Double free of p

```
void f() {  
    int* p = new int[100];  
    delete[] p;  
    delete[] p;  
}
```

## Q32. ⚡ What is wrong with deleting p twice?

- A) 📦 p becomes nullptr automatically
- B) 💀 Double free: undefined behavior
- C) ✂ Extra memory is freed safely
- D) 🚨 Compiler catches it at build time

```
std::vector<int> v;  
v.push_back(1);  
int* p = &v[0];  
v.push_back(2);  
*p = 99;
```

### Q33. What happens when you dereference 'p' after push\_back?

- A)  p still points to v[0]
- B)  Buffer overflow on the stack
- C)  Use-after-free: push\_back  
may reallocate the vector
- D)  Compile error: p is const

```
struct Foo { int x; Foo(){} } ;
Foo g_foo;
int main() {
    Foo s_foo;
    // g_foo.x vs s_foo.x ?
}
```

## Q34. What is g\_foo.x vs s\_foo.x before any assignment?

- A) g\_foo.x is undefined, s\_foo.x is 0
- B) Both are undefined
- C) Both are 0 (always zero-init)
- D) g\_foo.x is 0, s\_foo.x is undefined

## □ Memory

---

Stack: fast, auto, scope-bound (grows down)

Heap: dynamic, explicit new/delete/malloc/free

Global/Static: zero-init, lives for program lifetime

Stack bugs: overflow, dangling ref, buffer overrun

Heap bugs: leak, double-free, use-after-free

Modern C++: prefer RAII + smart pointers

```
$ ./my_tests  
--catch_system_errors=yes  
  
// What does this flag do?
```

## Q36. ● Boost.Test: what does -- **catch\_system\_errors=yes** do?

- A) Lets segfaults produce a core dump
- B) Only catches C++ exceptions
- C) Catches segfaults/signals and reports them as test failures
- D) Disables all error checking

```
// Adding a new feature...
// How to make it unit-testable?

class NewFeature {
    // ???
}
```

## Q37. 🧐 How do you make a new feature easy to unit test?

- A) Write the tests after the full feature is done
- B) Add it directly inside main()
- C) Implement it as a standalone class
- D) Use global variables for shared state

```
// Which framework has the  
// strongest mocks ecosystem?  
  
// A) Boost.Test  
// B) GoogleTest (gMock)  
// C) Catch2  
// D) All are equal
```

## Q38. 🤔 Which C++ test framework has the strongest mocks ecosystem?

- A) GoogleTest (gMock)
- B) Boost.Test
- C) Catch2
- D) All are equal



# Testing

---

Unit test: one class/function, fast (ms)

Component test: full algorithm + data replay

System test: whole project, threads, I/O, mocks

TDD: Red -> Green -> Refactor

Frameworks: GTest / Boost.Test / Catch2

CI: build -> test -> package, fail fast

```
int counter = 0;  
void run() {  
    for(int i=0;i<100000;++i)  
        counter++;  
}  
// Two threads call run(). Output?
```

## Q40. Two threads call run(). What is the output of counter?

- A) 0 (both threads cancel each other out)
- B) Unpredictable (data race, could be anything)
- C) A compile error
- D) Exactly 200000 (always)

```
std::condition_variable cv;
std::mutex mtx;
// Consumer:
std::unique_lock<std::mutex> lk(mtx);
cv.wait(lk); // bug?
process();
```

## Q41. 🧟 What is the bug in this `condition_variable` usage?

- A) Missing predicate: susceptible to spurious wakeups
- B) mtx must be a recursive\_mutex
- C) process() must be called before wait()
- D) cv.wait() needs a timeout argument

```
auto ptr = std::make_shared<int>(0);
// Thread A:
*ptr = 20;
// Thread B (concurrent):
*ptr = 30;
```

## Q42. Is writing \*ptr from two threads simultaneously safe?

- A) No: data race on the managed object
- B) Yes: shared\_ptr is fully thread-safe
- C) Yes: atomic ref-count protects the value
- D) Only unsafe on ARM

## □ Multithreading

---

Thread: OS resource sharing address space

Data race: two threads, one write -> UB

Mutex: use `lock_guard`, never `lock()` manually

CV: always use a predicate (spurious wakeups!)

`atomic<T>`: lock-free only for small types

Priority inversion: low holds lock, high blocks

```
double a = 0.1 + 0.2;  
if (a == 0.3)  
    std::cout << "Equal";  
else  
    std::cout << "Not Equal";
```

## Q44. 12 34 What does this code print?

- A) Equal (IEEE 754 is exact)
- B) Undefined behavior
- C) Not Equal (0.1+0.2 has rounding error)
- D) Compile error: == invalid for doubles

## Q45. What does **'double x = 1.0/0.0' print?**

```
double x = 1.0 / 0.0;  
std::cout << x;  
// What is printed?
```

- A) Crash: divide by zero
- B) 0
- C) +inf
- D) NaN

```
double n = std::sqrt(-1.0);
if (n == n)
    std::cout << "equal";
else
    std::cout << "not equal";
// What prints?
```

## Q46. 🤯 What does this NaN comparison print?

- A) equal (NaN is unique)
- B) Undefined behavior
- C) Compile error
- D) not equal (NaN != NaN is always true)

```
// Compiled with g++ -ffast-math
double a = 1e16;
double b = 1.0;
std::cout << (a + b) - a;
// Expected vs actual output?
```

## Q47. ⚡ What does `-ffast-math` do to floating-point code?

- A) Only speeds up integers, floats unchanged
- B) Allows non-IEEE optimizations, may change results
- C) Forces strict IEEE 754 compliance
- D) Disables SIMD instructions

## ☐ Numbers

---

IEEE 754: sign + exponent + mantissa

$0.1 + 0.2 \neq 0.3$  (finite binary precision)

`NaN != NaN` (always false, use `std::isnan`)

$1.0/0.0 = +\infty$ ,  $0.0/0.0 = \text{NaN}$  (no exception!)

`-ffast-math`: fast but non-IEEE, non-deterministic

ULP: smallest difference between two floats

```
std::tuple<int, std::string>
get_user_data() {
    return {1, "Alice"};
}

auto [id, name] = get_user_data();
```

## Q49. 17 What C++17 feature does 'auto [id, name] = ...' use?

- A) Structured bindings: unpack tuple/struct
- B) std::pair destructuring (C++11 feature)
- C) Template parameter pack
- D) Move semantics on return value

```
if (auto it = m.find("key");  
    it != m.end()) {  
    use(it->second);  
}  
// 'it' after the closing brace?
```

## Q50. 🔎 After the closing brace, is 'it' accessible?

- A) Yes: 'it' is a regular local variable
- B) No: 'it' is scoped to the if block
- C) Only if the if condition was true
- D) Only in the else branch

```
template<typename T>
void info(T val) {
    if constexpr (std::is_integral_v<T>)
        std::cout << val;
    else
        std::cout << val.size();
}
// info(42); -- would .size() compile?
```

## Q51. ⚙ Why does info(42) compile even though 'val.size()' is in the code?

- A) Template instantiation defers all checks
- B) The compiler ignores unreachable code paths
- C) int has a .size() method in C++17
- D) constexpr if discards the else branch at compile time for int

## 17[] C++17

---

Structured bindings: `auto [a, b] = tuple;`  
if/switch with initializer: `if (auto x=f(); x>0)`  
`std::optional<T>`: value or `std::nullopt`  
`std::filesystem`: portable path, `directory_iterator`  
`constexpr if`: compile-time branch (no SFINAE!)  
`std::string_view`: non-owning string reference